

# Message Chains for Distributed System Verification

[FEDERICO MORA](#), University of California, Berkeley

[ANKUSH DESAI](#), Amazon Web Services


[ELIZABETH POLGREEN](#), University of Edinburgh

[SANJIT A. SESHIA](#), University of California, Berkeley

“Distributed bugs... are often severe... Not only are these outages widespread and expensive, they can be caused by bugs that were deployed to production months earlier... [For example] [amazon.com] went down because one remote server couldn't display any product information.”


- Jacob Gabrielson, Senior Principal Engineer at Amazon Web Services

# Formal Verification To The Rescue!



Usually we must iterate,  
making updates to  $S$  or  $I$

# More Detailed $S$ ; More Complicated $I$



# Goal: Approach Implementation Level System Detail

```
receive  
  pong -> io:format("received pong~n", [])  
end,
```

Erlang

```
def receive = {  
  case "pong" => log.info("received pong")  
  case _    => log.info("received unknown message")  
}
```

Akka

```
on ePong do {  
  print "received pong";  
}
```

P

These are called  
event handlers

# Contributions Overview

- We formally define
  - message chains and
  - message chain invariants.
- We show how to use them to
  - verify asynchronous, message-passing systems
  - while avoiding
    - reasoning about individual messages or buffers,
    - manually abstracting message passing, and
    - manually relating abstractions to messaging details.
- Present the UPVerifier

# Running Example:

Chang And Roberts Algorithm For Decentralized  
Extrema-finding In Circular Configurations Of Processes

# Chang and Roberts Setting and Goal

Given a random circular arrangement of uniquely numbered processes where no a priori knowledge of the number of processes is known, and no central controller is assumed, we would like a method of designating by consensus a single unique process. The algorithm we



# Chang and Roberts Setting and Goal

Given a random circular arrangement of uniquely numbered processes

we would like a method of designating by consensus a single unique process.

# Chang and Roberts Algorithm

Each process is assumed to know its own number, and initially it generates a message with its own number, passing it to the left. A process receiving a message compares the number on the message with its own. If its own number is lower, the process passes the message (to its left). If its own number is higher, the process throws the message away, and if equal, it is the highest numbered process in the system.

# Chang and Roberts Algorithm

Each process is assumed to know its own number, and initially it generates a message with its own number, passing it to the left.

# Chang and Roberts Algorithm

A process receiving a message compares the number on the message with its own. If its own number is lower, the process passes the message (to its left).

# Chang and Roberts Algorithm

A process receiving a message compares the number on the message with its own.

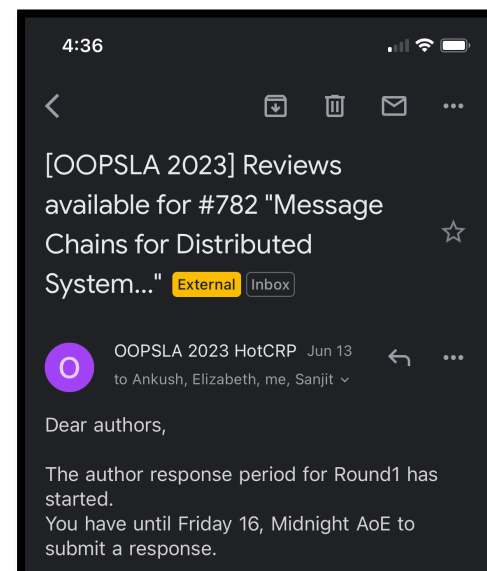
If its own number is higher, the process throws the message away

# Chang and Roberts Algorithm

A process receiving a message compares the number on the message with its own.

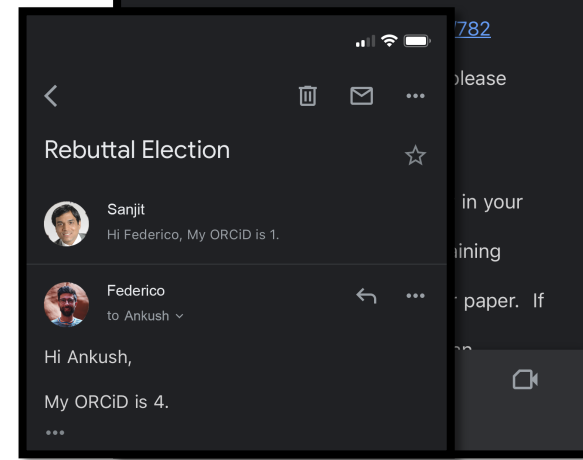
if equal, it is the highest numbered process in the system.

# Rebuttal Election Example



## Who Will Draft The Rebuttal?

1. ORCiD gives unique ID
2. Author list gives circle arrangement
3. Email gives message passing



# Example System Modelling:

Chang and Roberts Algorithm



# Chang and Roberts in TLA+

“msgs” is a global shared memory variable that they manually maintain

all duplicate messages dealt with at once

```
45   n1: while (TRUE) {
46       \* handle some incoming message
47       with (id \in msgs[self],
48           _msgs = [msgs EXCEPT ![self] = @ \ {id}]) {
49           if (state = "lost") { \* nodes that have already lost forward the message
               msgs := [_msgs EXCEPT ![succ(self)] = @ \cup {id}]
           } else if (id < Id[self]) {
               \* received smaller ID: record loss and forward the message
               state := "lost";
               msgs := [_msgs EXCEPT ![succ(self)] = @ \cup {id}]
           } else {
               \* do not forward the message; if it's the own ID, declare win
               msgs := _msgs;
               if (id = Id[self]) { state := "won" }
           }
       }
56   } \* end with
57
58
59
60
```

# Chang and Roberts in mpyvy

“pending” is a global shared memory variable that they manually maintain

```
28  transition recv(sender: node, n: node, next: node)
29      modifies leader, pending
30      (forall Z. n != next & ((Z != n & Z != next) -> btw(n,next,Z))) &
31      old(pending(sender, n)) &
32      (sender = n ->
33          (forall N. leader(N) <-> old(leader(N)) | N = n) &
34          (forall N1, N2.
35              !(N1 = sender & N2 = n) -> # message may or may not be removed
36              (pending(N1, N2) <-> old(pending(N1, N2))))) &
37      (sender != n ->
38          (forall N. leader(N) <-> old(leader(N))) &
39          (le(n, sender) ->
40              (forall N1, N2.
41                  !(N1 = sender & N2 = n) -> # message may or may not be removed
42                  (pending(N1, N2) <-> old(pending(N1, N2)) | N1 = sender & N2 = next))) &
43          (!le(n, sender) ->
44              (forall N1, N2.
45                  !(N1 = sender & N2 = n) -> # message may or may not be removed
46                  (pending(N1, N2) <-> old(pending(N1, N2)))))
```

simulate duplicates by possibly not removing message


# Other Choices, Same Story

- Domain Specific Tools: TLA+, mypyvy, etc...
- Proof Assistants: Lean, Coq, etc...
- Software Verification Engines: Dafny, F\*, etc...
- It is easier to write proofs about shared memory variables!

# Chang and Roberts in the UPVerifier

```
on eNominate e do {  
  let curr := e.payload.id in  
  if curr = this then  
    goto Won  
  else if this <= curr then  
    send left(this), eNominate(curr)  
}
```

# More Detailed $S$ ; More Complicated $I$



# Message Chains to the Rescue!

Message Chains Type,  
Message Chain Semantics,  
Message Chain Invariants,

# Message Chain Datatype

```
data [node_t, message_t] message_chain :=  
  | empty  
  | send {source: node_t,  
          target: node_t,  
          payload: message_t,  
          history: message_chain[node_t, message_t]}
```

# Verification Semantics of Send

## CASE 1/2: Send inside event handler

Receive a message chain

```
on eNominate e do {  
  let curr := e.payload.id in  
  if curr = this then  
    goto Won  
  else if this <= curr then  
    send left(this), eNominate(curr)  
}
```

Extend the received  
message chain

```
send(source = this,  
      target = left(this),  
      payload = eNominate(curr),  
      history = e)
```

the received message  
chain is the history of the  
new message chain



# Verification Semantics of Send

## CASE 2/2: Send outside event handler

No message received

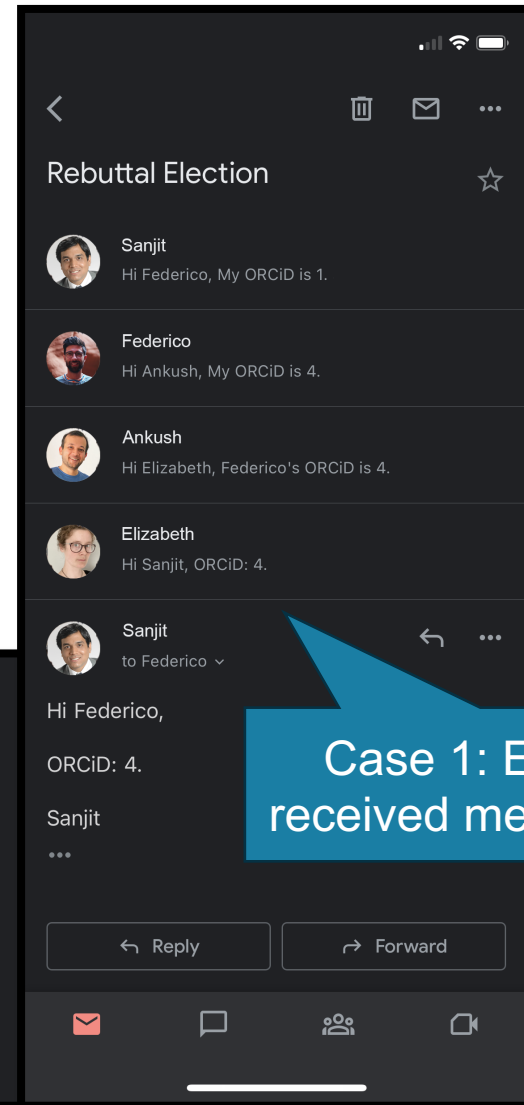
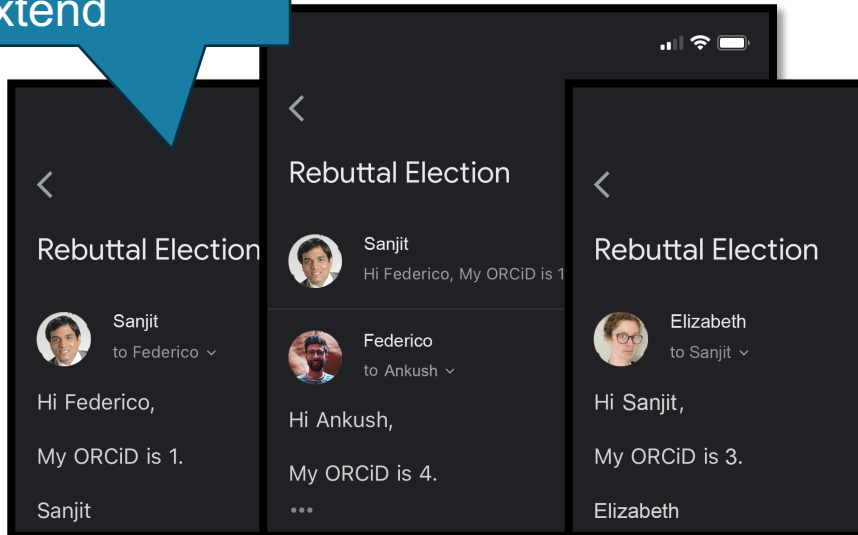
```
on entry do {  
    send left(this), eNominate(this)  
}
```

```
send(source = this,  
      target = left(this),  
      payload = eNominate(this),  
      history = empty)
```

There is no history to extend, so the history is empty

# Example Message Chains

Case 2: Nothing to extend



# Message Chain Invariants

s is the system, alive is true iff e is in a node's message buffer

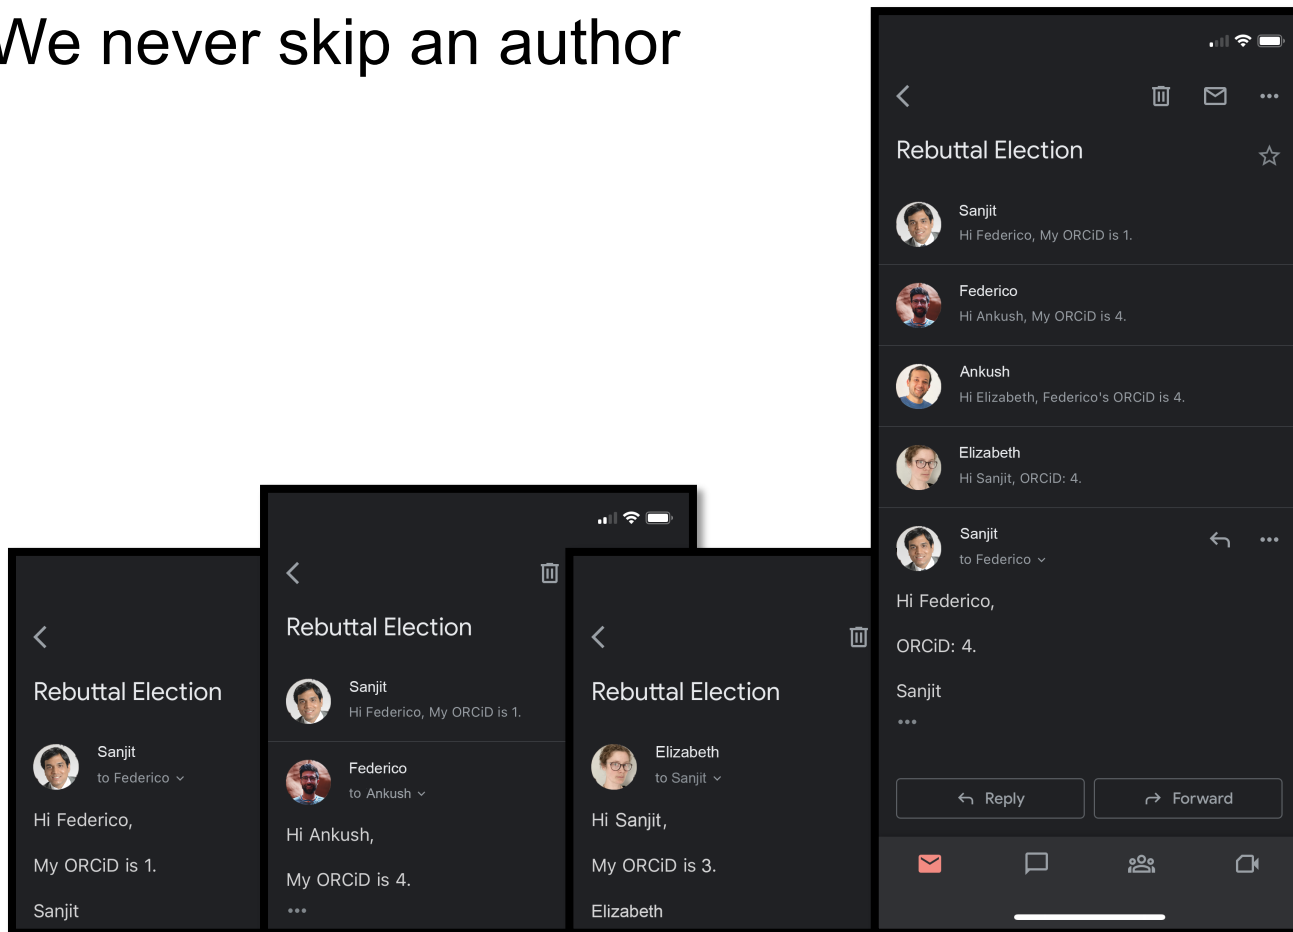
P is a predicate

forall (e: message\_chain[node\_t, message\_t]) alive(s, e) ==> P(e)

Everything is fixed except for P:  
P defines the invariant.

# Example Message Chain Invariants

1. Most recent email holds largest ORCiD
2. We never skip an author



# Example Verification With Message Chain Invariants:

Chang and Roberts Algorithm

# Chang and Roberts Full Verification

1. There is at most one leader:

```
forall (l: node, n: node) (leader(s, l) and leader(s, n)) ==> l = n
```

2. The leader has the largest ID:

```
forall (l: node, n: node) leader(s, l) ==> n <= l
```

## Two auxiliary message chain invariants: $P(e) :=$

1. The largest ID is always front:

```
forall (n: node) participated(e, n) ==> n <= e.payload.id
```

2. If a node has not participated, then the node is
  - the next target, or
  - the next target comes before that node in the ring:

```
forall (n: node) not participated(e, n) ==> between(e.payload.id, e.target, n) or e.target = n
```

Helper function that is true iff  $n$  is a source in  $e$

# Example Verification Analysis

- What does our proof guarantee?
  - At most one node will be elected no matter how many nodes are in the system.
- Why is our proof extra cool?
  - We never talk about individual messages or buffers,
  - invariants are (relatively) intuitive, and
  - abstraction is
    - created automatically and
    - automatically related to the implementation level detail

# Specification Mining



# Specification Mining Definition

- "... a machine learning approach to discovering formal specifications" that code likely obeys.
  - Input: observed program executions
  - Output: suggested invariant
- In our setting, we are interested in mining message chain invariants.
- Note: not auxiliary invariant inference, which requires a target specification.

# Specification Mining Example

```
[Empty] -> True,  
[Send(N1, N2, 1, Empty)] -> True,  
[Send(N2, N3, 2, Empty)] -> True,  
[Send(N1, N2, 4, Send(N4, N1, 4, Empty))] -> True,  
[Send(N4, N1, 3, Send(N3, N4, 2, Empty))] -> True,  
[Send(N4, N1, 4, Send(N3, N4, 3, Send(N2, N3, 2, Empty)))] -> True,  
[Send(N4, N1, 4, Send(N3, N4, 3, Send(N2, N3, 2, Send(N1, N2, 1, Empty))))] -> True,  
[Send(N3, N4, 4, Send(N2, N3, 4, Send(N1, N2, 4, Send(N4, N1, 4, Empty))))] -> True,
```

```
[Send(N2, N1, 0, Send(N1, N2, 1, Empty))] -> False,  
[Send(N1, N4, 3, Send(N3, N4, 4, Empty))] -> False,  
[Send(N4, N3, 4, Send(N4, N1, 5, Send(N4, N2, 2, Empty)))] -> False,
```

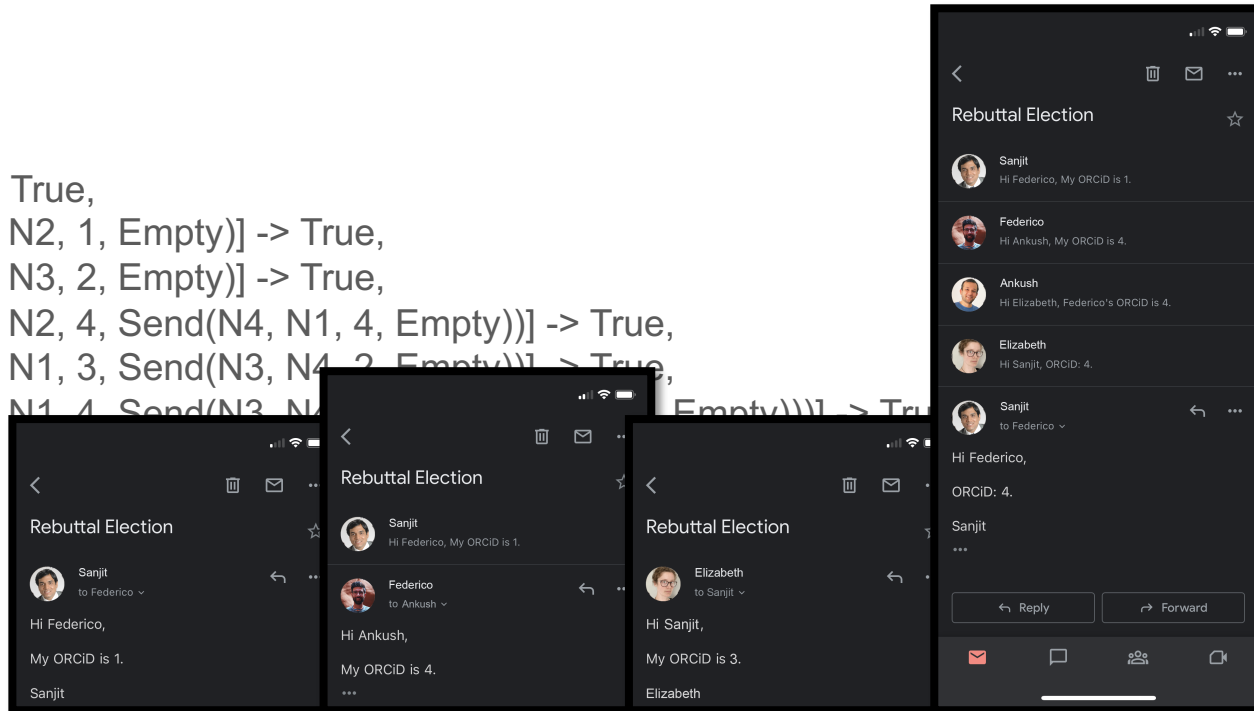
$P(e) := \text{forall } (n: \text{node}) \text{ participated}(e, n) \implies n \leq e.\text{payload.id}$

Actually, a recursive  
function that is  
equivalent to this

# Specification Mining Example

[Empty] -> True,  
 [Send(N1, N2, 1, Empty)] -> True,  
 [Send(N2, N3, 2, Empty)] -> True,  
 [Send(N1, N2, 4, Send(N4, N1, 4, Empty))] -> True,  
 [Send(N4, N1, 3, Send(N3, N4, 2, Empty))] -> True,  
 [Send(N4, N1, 4, Send(N3, N4, 2, Empty))] -> True,  
 [Send(N4, N1, 4, Send(N3, N4, 2, Empty))] -> True,  
 [Send(N3, N4, 2, Empty)] -> True,

[Send(N2, N3, 2, Empty)] -> True,  
 [Send(N1, N2, 4, Send(N4, N1, 4, Empty))] -> True,  
 [Send(N4, N1, 3, Send(N3, N4, 2, Empty))] -> True,



$P(e) := \text{forall } (n: \text{node}) \text{ participated}(e, n) \implies n \leq e.\text{payload.id}$

$P(e) :=$  "Most recent email holds largest ORCID"

# Empirical Evaluation

# Research Questions

- RQ1: Expressive power comparison to related work
  - Show we can translate proofs to the UPVerifier
- RQ2: Performance comparison to related work
- RQ3: Existing proofs with and without message chains
  - More case studies, like Chang and Roberts verification
- RQ4: Specification mining evaluation
  - More case studies, like Chang and Roberts spec. mining
- RQ5: Industrial distributed systems

# RQ2: Performance Comparison

Benchmark	mypyvy	UPVERIFIER
ring-id	0.365s	<b>0.138s</b>
toy-consensus-forall	0.380s	<b>0.064s</b>
consensus-wo-decide	0.338s	<b>0.072s</b>
sharded-kv	0.407s	<b>0.045s</b>
learning-switch	0.410s	<b>0.048s</b>
consensus-forall	0.566s	<b>0.120s</b>
lockserv	0.406s	<b>0.049s</b>
ticket	0.402s	<b>0.068s</b>
firewall	0.364s	<b>0.033s</b>
sharded-kv-no-lost-keys	<b>0.328s</b>	0.589s
client-server-ae	0.323s	<b>0.037s</b>
toy-consensus-epr	0.392s	<b>0.106s</b>
client-server-db-ae	0.403s	<b>0.060s</b>
ring-id-not-dead	0.479s	<b>0.194s</b>
consensus-epr	0.557s	<b>0.088s</b>
hybrid-reliable-broadcast	0.676s	<b>0.489s</b>

# RQ5: Industrial Distributed Systems

- GlobalClock
  - 150 lines of code.
  - ~1.5 seconds to verify.
  - Message chain invariant used to distinguish phases for local clocks.
- Two-Phase Commit Multi-Version Concurrency Control
  - 250 lines of code.
  - ~12 minutes to verify.
  - Interesting message chain invariants because of “broadcasting” and “flooding,” which are nonlinear.

# Summary and Thank you!

- We formally define
  - message chains and
  - message chain invariants.
- We show how to use them to
  - verify asynchronous, message-passing systems
  - while avoiding
    - reasoning about individual messages or buffers,
    - manually abstracting message passing, and
    - manually relating abstractions to messaging details.
- Present the UPVerifier



# Works Cited In Presentation

Ernest Chang and Rosemary Roberts. 1979. An improved algorithm for decentralized extrema-finding in circular configurations of processes. *Commun. ACM* 22, 5 (May 1979), 281–283. <https://doi.org/10.1145/359104.359108>

M. D. Ernst, J. Cockrell, W. G. Griswold and D. Notkin, "Dynamically discovering likely program invariants to support program evolution," in *IEEE Transactions on Software Engineering*, vol. 27, no. 2, pp. 99-123, Feb. 2001, doi: 10.1109/32.908957.

Glenn Ammons, Rastislav Bodík, and James R. Larus. 2002. Mining specifications. In *Proceedings of the 29th ACM SIGPLAN-SIGACT symposium on Principles of programming languages (POPL '02)*. Association for Computing Machinery, New York, NY, USA, 4–16. <https://doi.org/10.1145/503272.503275>

Woosuk Lee and Hangyeol Cho. 2023. Inductive Synthesis of Structurally Recursive Functional Programs from Non-recursive Expressions. *Proc. ACM Program. Lang.* 7, POPL, Article 70 (January 2023), 31 pages.

Jason R. Koenig, Oded Padon, Neil Immerman, and Alex Aiken. 2020. First-order quantified separators. In *Proceedings of the 41st ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI 2020)*. Association for Computing Machinery, New York, NY, USA, 703–717. <https://doi.org/10.1145/3385412.3386018>